

Kafka Serialization and Deserialization: JSON & Avro Examples

This is a corrected and updated guide for a Java client that demonstrates how to produce messages using **JSON** and **Avro** serialization. This version uses the latest compatible components to ensure a seamless experience.

1. Maven Project Configuration (pom.xml)

This pom.xml includes all the necessary dependencies for a project using Kafka clients, JSON serialization (via Jackson), and Avro serialization with the Confluent Schema Registry client. The versions have been updated to the latest stable releases.

```
<?xml version="1.0" encoding="UTF-8"?>
<project xmlns="http://maven.apache.org/POM/4.0.0"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://maven.apache.org/POM/4.0.0
http://maven.apache.org/xsd/maven-4.0.0.xsd">
  <modelVersion>4.0.0</modelVersion>

  <groupId>com.example</groupId>
  <artifactId>kafka-serialization-examples</artifactId>
  <version>1.0-SNAPSHOT</version>

  <properties>
    <maven.compiler.source>17</maven.compiler.source>
    <maven.compiler.target>17</maven.compiler.target>
    <project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
    <kafka.version>4.0.0</kafka.version>
    <jackson.version>2.16.1</jackson.version>
    <confluent.version>8.0.0</confluent.version>
    <avro.version>1.11.3</avro.version>
  </properties>

  <repositories>
    <repository>
      <id>confluent</id>
      <url>https://packages.confluent.io/maven/</url>
    </repository>
  </repositories>
```

```
<dependencies>
  <!-- Kafka Clients -->
  <dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-clients</artifactId>
    <version>${kafka.version}</version>
  </dependency>

  <!-- JSON Serialization with Jackson -->
  <dependency>
    <groupId>com.fasterxml.jackson.core</groupId>
    <artifactId>jackson-databind</artifactId>
    <version>${jackson.version}</version>
  </dependency>

  <!-- Avro Serialization with Confluent -->
  <dependency>
    <groupId>io.confluent</groupId>
    <artifactId>kafka-avro-serializer</artifactId>
    <version>${confluent.version}</version>
  </dependency>
  <dependency>
    <groupId>org.apache.avro</groupId>
    <artifactId>avro</artifactId>
    <version>${avro.version}</version>
  </dependency>

  <!-- SLF4J simple logger -->
  <dependency>
    <groupId>org.slf4j</groupId>
    <artifactId>slf4j-simple</artifactId>
    <version>1.7.36</version>
  </dependency>
</dependencies>
```

```
<build>
  <plugins>
    <!-- Maven Compiler Plugin -->
```

```

<plugin>
  <groupId>org.apache.maven.plugins</artifactId>
  <version>3.11.0</version>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <source>${maven.compiler.source}</source>
    <target>${maven.compiler.target}</target>
  </configuration>
</plugin>
<!-- Maven Assembly Plugin to create a fat JAR -->
<plugin>
  <groupId>org.apache.maven.plugins</artifactId>
  <artifactId>maven-assembly-plugin</artifactId>
  <version>3.6.0</version>
  <configuration>
    <descriptorRefs>
      <descriptorRef>jar-with-dependencies</descriptorRef>
    </descriptorRefs>
    <archive>
      <manifest>
        <mainClass>com.example.SerializationExamples</mainClass>
      </manifest>
    </archive>
  </configuration>
  <executions>
    <execution>
      <id>make-assembly</id>
      <phase>package</phase>
      <goals>
        <goal>single</goal>
      </goals>
    </execution>
  </executions>
</plugin>
</plugins>
</build>
</project>

```

2. Java Serialization Examples (SerializationExamples.java)

This class defines a User model and demonstrates a producer that sends messages as JSON strings and another that uses Avro's GenericRecord format, which is registered with a Schema Registry. The BOOTSTRAP_SERVERS and SCHEMA_REGISTRY_URL have been updated to use localhost for a more general example.

```
package com.example;

import com.fasterxml.jackson.databind.ObjectMapper;
import org.apache.avro.Schema;
import org.apache.avro.generic.GenericData;
import org.apache.avro.generic.GenericRecord;
import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerConfig;
import org.apache.kafka.clients.producer.ProducerRecord;
import org.apache.kafka.common.serialization.StringSerializer;
import io.confluent.kafka.serializers.KafkaAvroSerializer;

import java.io.IOException;
import java.util.Properties;

/**
 * A class demonstrating Kafka producers with JSON and Avro serialization.
 */
public class SerializationExamples {

    private static final String JSON_TOPIC = "json-topic";
    private static final String AVRO_TOPIC = "avro-topic";
    private static final String BOOTSTRAP_SERVERS = "localhost:9092";
    private static final String SCHEMA_REGISTRY_URL = "http://localhost:8081";

    public static void main(String[] args) throws InterruptedException {
        System.out.println("--- Running JSON Producer ---");
        runJsonProducer();
        Thread.sleep(2000);
    }
}
```

```

        System.out.println("\n--- Running Avro Producer ---");
        runAvroProducer();
        Thread.sleep(2000);
    }

    /**
     * An inner class to represent the data model for our examples.
     */
    public static class User {
        private String name;
        private int age;

        public User() {}

        public User(String name, int age) {
            this.name = name;
            this.age = age;
        }

        public String getName() { return name; }
        public void setName(String name) { this.name = name; }
        public int getAge() { return age; }
        public void setAge(int age) { this.age = age; }
    }

    /**
     * Creates and runs a producer that serializes a User object to a JSON string.
     */
    private static void runJsonProducer() {
        Properties props = createBaseProducerProperties();
        KafkaProducer<String, String> producer = new KafkaProducer<>(props);

        ObjectMapper objectMapper = new ObjectMapper();
        User user = new User("Alice", 30);

        try {
            String jsonString = objectMapper.writeValueAsString(user);
            ProducerRecord<String, String> record = new ProducerRecord<>(JSON_TOPIC,
"user-key-1", jsonString);

```

```

        producer.send(record, (metadata, exception) -> {
            if (exception == null) {
                System.out.println("JSON message sent successfully: " + jsonString);
            } else {
                exception.printStackTrace();
            }
        });
    } catch (IOException e) {
        e.printStackTrace();
    }

    producer.flush();
    producer.close();
}

/**
 * Creates and runs a producer that serializes a User object using Avro.
 * This requires the Schema Registry to be running.
 */
private static void runAvroProducer() {
    Properties props = createBaseProducerProperties();
    // The value serializer for Avro messages
    props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
KafkaAvroSerializer.class.getName());
    // The URL of the Schema Registry
    props.put("schema.registry.url", SCHEMA_REGISTRY_URL);

    KafkaProducer<String, GenericRecord> producer = new KafkaProducer<>(props);

    // Define the Avro schema inline
    String userSchemaString =
"{\"type\": \"record\", \"name\": \"User\", \"fields\": [{\"name\": \"name\", \"type\": \"string\"}, {\"name\": \"age\", \"type\": \"int\"}]}"
    Schema schema = new Schema.Parser().parse(userSchemaString);

    // Create a GenericRecord from the schema
    GenericRecord userRecord = new GenericData.Record(schema);
    userRecord.put("name", "Bob");
    userRecord.put("age", 45);

```

```

        // Send the Avro message
        ProducerRecord<String, GenericRecord> record = new
        ProducerRecord<>(AVRO_TOPIC, "user-key-2", userRecord);
        producer.send(record, (metadata, exception) -> {
            if (exception == null) {
                System.out.println("Avro message sent successfully: " + userRecord);
            } else {
                System.err.println("Error sending Avro message: " +
                exception.getMessage());
            }
        });

        producer.flush();
        producer.close();
    }

    /**
     * Common producer properties.
     */
    private static Properties createBaseProducerProperties() {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
        BOOTSTRAP_SERVERS);
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
        StringSerializer.class.getName());
        // For the JSON producer, the value serializer is a StringSerializer
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
        StringSerializer.class.getName());
        return props;
    }
}

```

3. Step-by-Step Execution Guide

Follow these steps to set up the project, run the necessary services, and execute the examples.

Step 1: Project Setup and Compilation

1. Create a project directory structure: `mkdir -p kafka-serialization-examples/src/main/java/com/example`.
2. Copy the `pom.xml` and `SerializationExamples.java` content into the correct locations.
3. Navigate to the project root and build the fat JAR: `mvn clean package`.

Step 2: Install Confluent Platform

This is an essential step for the Avro example. The example assumes you have a user named `kafka` and your Kafka is installed in `/usr/local/kafka`.

```
# Define versions for Confluent and Kafka
```

```
CONFLUENT_VERSION_DL="8.0.0"
```

```
CONFLUENT_VERSION="8.0"
```

```
KAFKA_VERSION="4.0.0"
```

```
# Download the Confluent Platform tarball
```

```
wget
```

```
"https://packages.confluent.io/archive/${CONFLUENT_VERSION}/confluent-community-${CONFLUENT_VERSION_DL}.tar.gz" -O /tmp/confluent.tar.gz
```

```
# Extract to a new directory
```

```
sudo mkdir -p /usr/local/confluent
```

```
sudo tar -xzf /tmp/confluent.tar.gz -C /usr/local/confluent --strip-components 1
```

```
# Change ownership to the kafka user
```

```
sudo chown -R kafka:kafka /usr/local/confluent
```

Note: The community download filename for Confluent Platform 8.0.0 is slightly different, removing the Kafka version suffix.

Step 3: Start Confluent Schema Registry

With the Confluent Platform installed, you can now start the Schema Registry.

Update

```
sudo vim /usr/local/confluent/etc/schema-registry/schema-registry.properties
```

```
#listeners=http://0.0.0.0:8081
```

```
listeners=http://kafka-199:8081
```

```
# Use this setting to specify the bootstrap servers for your Kafka cluster and it
```

```
# will be used both for selecting the leader schema registry instance and for storing  
the data for
```

```
# registered schemas.
```

```
#kafkastore.bootstrap.servers=PLAINTEXT://localhost:9092
```

```
kafkastore.bootstrap.servers=PLAINTEXT://kafka-199:9092,kafka-200:9092,kafka-201  
:9092
```

```
# The name of the topic to store schemas in
```

```
kafkastore.topic=_schemas
```

```
# If true, API requests that fail will include extra debugging information, including  
stack traces
```

```
debug=false
```

```
metadata.encoder.secret=REPLACE_ME_WITH_HIGH_ENTROPY_STRING
```

```
resource.extension.class=io.confluent.dekregistry.DekRegistryResourceExtension
```

```
sudo -u kafka /usr/local/confluent/bin/schema-registry-start  
/usr/local/confluent/etc/schema-registry/schema-registry.properties
```

Step 4: Create the Topics

Create the two topics that the producers will write to.

```
# JSON Topic
```

```
sudo -u kafka /usr/local/kafka/bin/kafka-topics.sh --create \  
--topic json-topic \  

```

```
--partitions 3 \  
--replication-factor 3 \  
--bootstrap-server localhost:9092
```

Avro Topic

```
sudo -u kafka /usr/local/kafka/bin/kafka-topics.sh --create \  
--topic avro-topic \  
--partitions 3 \  
--replication-factor 3 \  
--bootstrap-server localhost:9092
```

Step 5: Run the Java Producer Application

Execute the JAR file to run both producer examples sequentially.

```
java -jar  
target/kafka-serialization-examples-1.0-SNAPSHOT-jar-with-dependencies.jar
```

Step 6: Verify Messages with a Consumer

Open two separate terminals to verify the messages produced to each topic.

Consumer for JSON Messages:

Since the JSON message is just a string, you can use the default consumer.

```
sudo -u kafka /usr/local/kafka/bin/kafka-console-consumer.sh \  
--topic json-topic \  
--bootstrap-server kafka-199:9092 \  
--from-beginning \  
--property print.key=true
```

You should see: user-key-1 {"name":"Alice","age":30}

Consumer for Avro Messages:

To read Avro messages, you must use the Confluent-provided consumer script and specify the Schema Registry URL. The kafka-avro-console-consumer is the correct tool to use for this purpose.

```
sudo -u kafka /usr/local/confluent/bin/kafka-avro-console-consumer \  
  --topic avro-topic \  
  --bootstrap-server kafka-199:9092 \  
  --from-beginning \  
  --property print.key=true \  
  --property  
key.deserializer=org.apache.kafka.common.serialization.StringDeserializer \  
  --property schema.registry.url=http://kafka-199:8081
```

You should see: user-key-2 {"name": "Bob", "age": 45}